


BUILD


//LOCALHOST:VIENNA

Tuesday, June 16, 2026 · 14:00–18:00
Microsoft Austria



Agenda

Time	Session	Speaker
14:00	Welcome	Bernd Zimmermann, Microsoft
14:05	Build Recap, Claws agent harness in Microsoft Foundry	Martina Grom, MVP Toni Pohl, MVP
14:30	Build context-aware agents that scale: From data to decisions	Jürgen Mayrbäurl, MVP
15:15	Designing Applications with Microsoft IQ and AI-ready Databases	Peter Paul Kirschner, MVP
15:45	Break	
16:15	From CLI to PR: Automating the path to merged code	Patrick Wahlmüller, MVP
16:45	How GitHub Copilot can help development with Visual Studio 2026	Christian Nagel, MVP
17:30	Agentic AI that we actually use - self built and not ordered	Georg Binder, Microsoft
18:00	End	




BUILD Recap

The Agent Era Arrives

martina.grom@atwork.at · MVP Foundry and M365

toni.pohl@atwork.at · MVP Foundry and M365 Dev



Microsoft BUILD 2026

Keynote

NEWS

100+ announcements

IMPACT

The full agentic computing stack from silicon to models to apps

RESOURCES

<https://developer.microsoft.com/blog/build-recap>

“AI alone won’t change your business. The system running it will.” – KEY01

Satya Nadella, Build 2026 Keynote

SYSTEM

Microsoft CEO Satya Nadella opened Build 2026 declaring AI models alone aren’t enough - you need the entire system running them, defined as six interconnected layers.

HOW

The six-layer system: Compute → Models → Context (IQ) → Tools → Runtime → Security/Governance - each layer enabling the next, from hardware up to the Teams or Windows surface.

WHY IT MATTERS

This framing explains every Build 2026 announcement and is the architecture map every developer building production AI applications needs to keep in their head.



500+
datacenters



80+
Azure regions

you sort of have the data
the -- the layer about

Security & observability

Developer tools

Agent runtime

Models • Context • Tools

Infrastructure





Agenda

30 slides across 8 domains - from silicon to models to apps

MODELS

Seven New MAI Models · Frontier Tuning
Aion 1.0 - On-Device SLMs for Windows

AGENTS & FOUNDRY

Microsoft Foundry (12,000+ models) · Microsoft IQ
Microsoft Agent Framework 1.0 · Microsoft Scout

GOVERNANCE

Agent Control Spec (ACS) · Agent 365 SDK
MXC · Windows 365 for Agents · Copilot Credits

GITHUB & COPILOT

GitHub Copilot App (Agent-native desktop)
Copilot SDK GA · MAI-Code-1-Flash · Code Review

WINDOWS

Windows as the Agent OS · WSL Containers
Intelligent Terminal · Coreutils & Dev Configurations

SECURITY

MDASH - Agentic Scanning (96.55% CyberGym)
Defender + GitHub Code Security (GA)

AZURE

Azure HorizonDB · Cobalt 200 (+50% perf)
Microsoft Discovery (GA) · OpenClaw on Windows

HARDWARE

Project Solara · Majorana 2 Quantum
Surface RTX Spark Dev Box (1 petaflop)

7 New Microsoft AI Models

MAI
Image-2.5



MAI
Image-2.5-Flash



MAI
Transcribe-1.5



MAI
Thinking-1



MAI
Voice-2



MAI
Voice-2-Flash



MAI
Code-1-Flash



Seven New MAI Models – DEM323

Microsoft's biggest first-party model release - zero OpenAI distillation

WHAT IT IS

Microsoft AI Superintelligence Team (led by Mustafa Suleyman) launched seven first-party models built from scratch with zero OpenAI distillation - the biggest in-house model release in Microsoft's history.

WHAT IT DOES

Lineup: MAI-Thinking-1 (35B reasoning), MAI-Image-2.5 (image gen), MAI-Transcribe-1.5 (speech), MAI-Voice-2 (voice), MAI-Code-1-Flash (coding), plus Mayo Clinic health frontier models - available on Azure Foundry, Fireworks AI, Baseten, and OpenRouter.

WHY IT MATTERS

Developers get genuine first-party model choice competitive with leading third-party options, and enterprises gain an alternative AI supply chain independent of any single vendor.

Seven New MAI Models

Microsoft’s biggest first-party model release - zero OpenAI distillation

WHAT IT IS

DEM323

Under the hood of Microsoft AI models

#	Model	Category
1	MAI-Thinking-1	Reasoning - 35B parameter model for complex enterprise and coding tasks
2	MAI-Code-1-Flash	Coding - purpose-built for GitHub Copilot, small-tier, best-in-class for its size
3	MAI-Image-2.5	Image generation - full-quality image model
4	MAI-Image-2.5-Flash	Image generation - faster/lighter variant
5	MAI-Transcribe-1.5	Multilingual speech-to-text transcription
6	MAI-Voice-2	Voice synthesis and voice tasks
7	MAI-Voice-2-Flash	Voice - faster/lighter variant

Private preview

Frontier Tuning

Tuned workflow
environments

Continuous model
optimization

Enterprise-safe
by design

aka.ms/FrontierTuning

Frontier Tuning – DEM323, DEM321

Make the model fit your company - hill-climbing optimization on your own data

WHAT IT IS

Frontier Tuning is Microsoft's enterprise model customization system: a “hill-climbing” optimization loop that continuously tests and refines a MAI model against a company's own data and evaluation criteria until performance peaks.

WHAT IT DOES

The Land O'Lakes demo showed Frontier Tuning making a MAI model outperform generic models on domain-specific tasks (butter quality reporting), running automatically without sending proprietary data to a third party.

WHY IT MATTERS

Developers building on MAI models can now fine-tune them on proprietary data with no cloud transfer - closing the gap between frontier capability and enterprise-specific accuracy.

Aion 1.0 - On-Device SLMs for Windows

Always available · always private · works even offline

WHAT IT IS

BRK260

Build Apps w/ Local AI
for Unmetered Intelligence
on every Windows PC

The image shows a presentation slide for 'Aion 1.0 Instruct' with a presenter standing on a stage in front of the screen. The slide features the Windows logo at the top center, followed by the title 'Aion 1.0 Instruct'. Below the title, three key performance metrics are listed: '3.4x smaller memory footprint', '6x faster summarization', and '2x faster responses'. The URL 'aka.ms/Aion1' is displayed twice: once in the center of the slide and once in a white box at the bottom left of the stage.

Announcing



Aion 1.0 Instruct

3.4x smaller
memory footprint

6x faster
summarization

2x faster
responses

aka.ms/Aion1

aka.ms/Aion1

Aion 1.0 - On-Device SLMs for Windows – BRK260

Always available · always private · works even offline

WHAT IT IS

Aion 1.0 Instruct is a fast on-device SLM for summarization, rewrites, intent detection, and accessibility - running on the Windows device without a cloud round-trip. Aion 1.0 Plan is a 14B-parameter reasoning and tool-calling model with 32K context, shipping in-box with Windows.

WHAT IT DOES

Together they enable fully local agentic workflows: agents that reason, plan, and execute entirely on a Windows 11 PC with zero network dependency.

WHY IT MATTERS

Developers can build always-available, privacy-preserving, offline-capable agents on any Windows 11 PC - making agentic experiences viable in air-gapped, regulated, or connectivity-constrained environments.

...ary, so we're
to bring the frontier

... continuing t



Foundry Models

11,000+ models



Microsoft Foundry – BRK243, ODSP934, DEMSP383

Production platform for the full agent lifecycle - 12,000+ models

WHAT IT IS

Azure AI Foundry is Microsoft's production platform for the full agent lifecycle - authoring, evaluation, deployment, monitoring, and improvement - expanded to 11,000+ models including OpenAI GPT-5.5, Anthropic Claude Opus 4.8 (preview), Fireworks AI (GA), and all seven MAI models.

WHAT IT DOES

Hosted Agents (GA by end of June) adds hypervisor-isolated sandboxing, per-agent Entra ID, source-code deployment via azd, built-in content safety, and Voice Live + WebSocket support, wrapping every agent in security and identity from the start.

WHY IT MATTERS

Developers get a single platform to build, run, evaluate, and improve agents at production scale without managing infrastructure - replacing a patchwork of tools with one coherent lifecycle system.



aka.ms/MicrosoftIQ

Microsoft IQ – DEM323, LIVE171

Work IQ · Web IQ · Foundry IQ - grounding agents in real knowledge

WHAT IT IS

Microsoft IQ is the unified context layer that grounds agents in real knowledge: Work IQ connects to Microsoft 365 data (Teams, Outlook, SharePoint, OneDrive); Foundry IQ is a serverless enterprise knowledge retrieval layer; Web IQ provides live grounding from the fresh open web.

WHAT IT DOES

All three are generally available and accessible from GitHub Copilot, Azure Foundry, and Copilot Studio - enabling agents to pull current, contextual knowledge without building custom RAG pipelines.

WHY IT MATTERS

For the first time developers have a production-grade grounding layer spanning from the live web to the company intranet - no custom connectors or bespoke vector databases required.

Microsoft Agent Framework 1.0 – BRK243, DEM361, OD840

Open foundation for building AI agents - GA in Python & .NET

WHAT IT IS

Agent Framework (v1.0) → Build and run autonomous agents.


Microsoft Agent Framework 1.0 is now generally available in Python and .NET - a platform-agnostic library for authoring, wiring, and running AI agents that works across any model or AI service, not just Microsoft's.

WHAT IT DOES

The framework includes multi-agent orchestration, procedural memory, adaptive control, and a new Agent Control Specification (ACS) for runtime governance - the same foundation powering Copilot Studio, Copilot, and OpenClaw.

WHY IT MATTERS

Building on Agent Framework 1.0 means investing in Microsoft's long-term platform commitment - the same open SDK that ships in Microsoft's flagship enterprise AI products.



Demo

Preview



Secure local agents and claws with Agent 365

Discover and
secure local agents

Control agent access with
Entra + Intune policies


Protect data and defend against
threats with Purview + Defender

Run agents safely in
Windows sandbox, MXC

aka.ms/A365Local

aka.ms/A365Local

Agent 365 - BRK251, OD840



SDK GA · Agent Registry · Local Agent Discovery

WHAT IT IS

Agent 365 SDK is now generally available - free and framework-agnostic - embedding observability, access controls, and compliance enforcement directly into how agents are designed rather than bolted on after they misbehave.

WHAT IT DOES

The Agent 365 Agent Registry uses Defender, Entra, and Intune to surface unmanaged local agents an organization didn't know were running - recognizing 20+ agent types including coding agents and MCP servers.

WHY IT MATTERS

This shifts enterprise AI governance from a post-incident activity to a pre-deployment gate - security and IT teams gain visibility and control over the entire agent fleet through familiar tools.

Microsoft Scout – KEY01

The first Autopilot - always-on autonomous agent with its own identity

WHAT IT IS

Microsoft Scout is the first “Autopilot” - a new category of always-on autonomous agent with its own Entra ID identity, available as a desktop app for Win 11+ / macOS 12+, and the first AI agent designed to take initiative in enterprise environments rather than just responding to prompts.

WHAT IT DOES

Built on OpenClaw and powered by Work IQ, Scout accesses Teams, Outlook, OneDrive, and SharePoint to coordinate projects, surface risks, prepare meeting summaries, and resolve scheduling conflicts - available today in the Frontier program preview.

WHY IT MATTERS

Scout signals the end of the reactive AI assistant era - persistent agents acting as delegated workers with their own identity, audit trail, and enterprise governance are now a production reality.

OpenClaw on Windows – LIVE144, BRK245

Open-source agent framework - native on Windows with OS-level security

WHAT IT IS

OpenClaw is Microsoft's open-source agent framework for building autonomous, composable agents ("claws") that chain tasks, call tools, and collaborate across local and cloud environments - now open-source with a native Windows companion app.


WHAT IT DOES

At Build 2026, OpenClaw runs inside Microsoft Execution Containers (MXC) for security isolation, is the runtime foundation powering Microsoft Scout, and NVIDIA is integrating it through OpenShell on Windows.

WHY IT MATTERS

By open-sourcing the same agent substrate that powers Scout, Microsoft invites developers to build on production-proven infrastructure - not a toy framework, but the one shipping in Microsoft's most ambitious enterprise AI product.

Microsoft Execution Containers



Microsoft Execution Containers (MXC) – BRK262, BRKSP92

OS-level policy enforcement - Windows decides what agents are allowed to do

WHAT IT IS

The Microsoft Execution Container SDK is a new OS-level policy engine letting developers declare exactly what an agent can access - files, network, processes - with containment boundaries enforced by Windows at runtime, not by the agent itself.


WHAT IT DOES

MXC offers a composable spectrum of isolation semantics (lightweight sandbox to full VM), integrates with Defender, Entra, Intune, and Purview, and is the layer making OpenClaw, Scout, and NVIDIA OpenShell safe to deploy on enterprise PCs.

WHY IT MATTERS

The Build demo - an agent failing to delete a folder stopped by MXC policy - illustrated the correct mental model: the future of Windows agents is Windows deciding what the agent can do, shifting trust from the agent to the platform.


Agent Execution Models: From Open Access to Full Isolation



No sandbox

YESTERDAY


The agent runs directly on the user's machine with full access to the file system, network, and installed tools.



Sandboxed Tools

TODAY

The main agent loop runs on the host with low permissions, but delegates tool execution to isolated containers.



Fully isolated

TOMORROW

Agents executing within the sandbox have minimal to no knowledge of the resources and identities being used.

Windows 365 for Agents – OD855, OD852, DEM354

Cloud PCs for agent workloads - isolated, governed, auditable

WHAT IT IS

Windows 365 for Agents is now generally available, giving computer-using agents their own dedicated, policy-governed Cloud PC to run enterprise workflows in a secure, isolated environment - treating agents as managed compute tenants, not shadow processes.

WHAT IT DOES

Each agent Cloud PC has its own Entra ID, applies Defender and Intune policies automatically, and can be provisioned and deprovisioned programmatically via API - giving IT the same lifecycle management for agents as for human employees.

WHY IT MATTERS

Developers building enterprise automation get a real identity, a bounded execution environment, and a full audit trail for every agent - making regulatory compliance achievable for autonomous workloads.

Announcing



GitHub Copilot app

Start your agent sessions
from your GitHub workflow

Run parallel sessions in
isolated worktrees & branches

Drive PRs through review
to merge – all in one app

aka.ms/GitHubCopilotApp

aka.ms/GithubCopilotApp

GitHub Copilot App – LIVE152

Agent-native desktop experience - direct multiple coding agents in parallel

WHAT IT IS

The GitHub Copilot App is a new agent-native desktop experience (technical preview) with a single “My Work” view showing active agent sessions, issues, PRs, and automations across repos - available for Copilot Pro, Pro+, Business, and Enterprise users.

WHAT IT DOES

Key capabilities: parallel agent sessions in isolated git worktrees, bidirectional Canvases where agents update diffs and plans in real time, Agent Merge that carries PRs through CI and review automatically, and cloud/local sandboxes for safe code execution.

WHY IT MATTERS

This is the biggest architectural shift in GitHub Copilot’s history - from a pair programmer that responds to prompts to an agent orchestration platform where developers set direction and multiple agents execute work in parallel.

Copilot SDK & Copilot CLI - BRK206, LIVE152

GA in 6 languages · redesigned CLI with voice & scheduled tasks

WHAT IT IS

The Copilot SDK lets you create custom agents: Extend Copilot experiences and integrate agents into apps.

The Copilot SDK is now generally available in six languages - Node.js/TypeScript, Python, Go, .NET, Rust, and Java - exposing the same agentic runtime powering the Copilot App to any developer.

WHAT IT DOES

The Copilot CLI received a full redesign with a new TUI, on-device voice input (audio never leaves your machine), and a /every command for scheduling recurring automated tasks.

WHY IT MATTERS

With GA SDK in the top six developer languages and a scriptable, voice-enabled CLI with scheduling, Copilot is now a platform any developer can build on in their language of choice - not a feature locked inside VS Code.

presentation-2026-06-Build-CopilotSDK | ● 10 |   Update     - 

     ... 

AppState.cs X # Agent.cs TS chatSocket.ts TS environment.ts

demos > AgentOrchestrator > # demos\AgentOrchestrator\AppState.cs > { } AgentOrchestrator >  AppState >  copilotClient : CopilotCli

1 using GitHub.Copilot;
2
3 namespace AgentOrchestrator;
4
5 public class AppState(PropertyDatabase propertyDatabase)
6 {
7 private readonly CopilotClient copilotClient = new CopilotClient(new()
8 {
9 Mode = CopilotClientMode.Empty,
10 BaseDirectory = Path.Combine(AppContext.BaseDirectory, ".copilot"),
11 });
12
13 public Dictionary<string, Agent> Agents { get; } = new();
14 public event Action? UpdateUi;
15
16 public async Task RunAgentAsync(string enquiry)
17 {

Windows as the Agent OS – DEM345, OD858, DEMSP384

Windows Agent Framework - AI agents as first-class Windows citizens

WHAT IT IS

Windows Agent Framework introduces OS-level APIs letting AI agents register as first-class Windows services - appearing in the taskbar, scheduling tasks, receiving file system events, and managed through Intune - rather than running as unmanaged overlay applications.

WHAT IT DOES

The Windows Agent Runtime (Insider preview, June 2026) is the execution layer; it starts text-first with multi-modal agents following in subsequent releases as security validation catches up to capability.

WHY IT MATTERS

Microsoft is evolving Windows from a desktop for humans to an execution environment for both humans and autonomous software - with 1.4 billion active Windows devices, the scale implications of this shift are enormous.

WSL Containers & Intelligent Terminal – DEM346, BRK261

Built-in Linux containers + AI-native terminal - no Docker Desktop required

WHAT IT IS

WSL containers brings built-in Linux container creation and management to Windows using familiar CLI and APIs - no Docker Desktop required - with WSL as the container runtime, coming soon to public preview.


WHAT IT DOES

The Intelligent Terminal is an experimental Windows Terminal upgrade embedding GitHub Copilot context awareness and ACP integration directly in the terminal, letting agents debug errors and run multi-step tasks while the developer stays in their shell.

WHY IT MATTERS


For developers who live in the terminal, these two features make Windows feel like a native Linux/macOS environment: Linux containers without Docker Desktop, and a terminal with agent intelligence built in rather than a separate chat window.

We combine **multiple models, 100s of agents**, and unmatched **security intelligence** to find **real risk** – not theoretical noise





Run



Open



My work

Automations

Search

Sessions



Quick chats

mdash-corp

Defender fix



Defender upload

Defender details

Defender scan

BigM-Project

Project-Ni-Ni

Team-Grayson



Sarah Young



DEFENDER SCAN — 63 findings

● CRITICAL (3)

- CWE-78: OS Command Injection — process_data()
routes/public_api.py:278 → Confidence: LIKELY
- CWE-502: Deserialization — configure_template()
routes/reports.py:167 → Confidence: LIKELY
- CWE-208: Observable Timing Discrepancy — verify_password()
models/auth.py:110 → Confidence: LIKELY

● HIGH (5)

- CWE-1336: Improper Neutralization of Special Elements Used in Template Engine — render_page()
routes/public_api.py:39 → Confidence: LIKELY
- CWE-918: Server-Side Request Forgery — webhook_test()
routes/public_api.py:211 → Confidence: LIKELY
- CWE-1336: Improper Neutralization of Special Elements ' d in Template Engine — custom_report()
routes/reports.py:98 → Confidence: LIKELY

Ask anything. Type / for commands, @ to add files, or # to reference issues...



Interactive



GPT-5.5



Medium



MDASH - Agentic Security Scanning – BRK251, OD841

100+ AI agents · 96.55% CyberGym benchmark · up ~10 points in 3 weeks

WHAT IT IS

MDASH (Microsoft Security Multi-model Agentic Scanning Harness) orchestrates 100+ specialized AI agents across a configurable multi-model ensemble to discover, validate, and prove exploitability across codebases - achieving a CyberGym benchmark score of 96.55%.

WHAT IT DOES

MDASH integrates with Microsoft Defender for production signals (internet exposure, data sensitivity, blast radius) and routes AI-generated fix suggestions through GitHub Copilot as suggested PRs - so developers see the vulnerability, proof of exploitability, and a fix in one workflow.

WHY IT MATTERS

For security teams drowning in theoretical scan findings, MDASH represents a structural shift - proving actual exploitability before raising an alert, moving from “potential issues” to “confirmed exploitable bug with fix ready for review.”




Defender fix · sarah-yo/defender-remediation-fix



Run



Open



Ready to merge



My work

Automations

Search

Sessions



Quick chats

mdash-corp

Defender fix

Defender upload

Defender details

Defender scan

BigM-Project

Project-Ni-Ni

Team-Grayson



Sarah Young



Read changes 1 commit, 14 files

Edit reports.py +1 -1

Run final checks and diff stat

\$ Set-Location 'C:\Users\sayoung\copilot-worktrees\mdash-corp\sarah-yo-curly-spork'

python -m ruff check .

if (\$LASTEXITCODE -ne 0) { exit \$LASTEXITCODE }

python -m pytest -q

git --no-pager diff --stat

models/order.py | 24 ++++++

routes/admin.py | 4 +--

routes/coupons.py | 12 +++++--

routes/feedback.py | 30 ++++++-----

routes/files.py | 49 ++++++-----

routes/imports.py | 32 ++++++----- ↓

routes/orders.py | 2 ++

Changes +421 -137 PR #3

Ask anything. Type / for commands, @ to add files, or # to reference issues...

+ Interactive

GPT-5.5

Medium



Security and quality

Overview

Findings

Code quality

Dependabot

Malware

Vulnerabilities

Code scanning

89

Reporting

Security policy

Advisories

Code scanning

✓ All tools are working as expected

Tools 2 + Add tool

is:open branch:main tool:defender-aiscan

✕ Clear current search query, filters, and sorts

63 Open 0 Closed

Language Tool Rule Severity Sort

☐ ⚠ **Deserialization of Untrusted Data** ⚠ Warning
#89 opened yesterday • Detected by defender-aiscan in routes/reports.py :168

☐ ⚠ **XML External Entity Reference** ⚠ Warning
#88 opened yesterday • Detected by defender-aiscan in routes/imports.py :29

☐ ⚠ **Deserialization of Untrusted Data** ⚠ Warning
#87 opened yesterday • Detected by defender-aiscan in routes/public_api.py :104

☐ ⚠ **Improper Neutralization of Special Elements Used in Template Engine** ⚠ Warning
#86 opened yesterday • Detected by defender-aiscan in routes/search.py :174

☐ ⚠ **Eval Injection** ⚠ Warning
#85 opened yesterday • Detected by defender-aiscan in routes/public_api.py :59

main

main

main

main

main

Defender + GitHub Code Security Integration – BRK251, OD841

Vulnerability discovery enriched with production signals - fix in one workflow

WHAT IT IS

The integration between Microsoft Defender and GitHub Code Security (formerly GitHub Advanced Security) is now generally available, enriching every vulnerability with production signals - internet exposure, data sensitivity, blast radius - before surfacing it to developers.


WHAT IT DOES

AI-generated fixes are routed through GitHub Copilot as suggested pull requests, so developers see the vulnerability in context, the proof of exploitability from Defender, and a proposed code fix in a single PR workflow.

WHY IT MATTERS

This bridges the persistent gap between security teams who find issues and developers who fix them - security remediation becomes part of the standard PR workflow rather than a competing queue that slows feature work.

More stuff...



Project Solara

Chip-to-cloud platform for ambient, always-on AI agent hardware

WHAT IT IS

Project Solara is Microsoft's new chip-to-cloud platform for agent-first hardware: a Qualcomm-powered AI badge with biometric auth and camera, a smart desk companion running Windows 365, and concept designs for glasses, watches, speakers, and earbuds.

WHAT IT DOES

The platform provides reference architectures and SDKs for agent-first devices so developers can build autonomous agent experiences across any Solara-compatible form factor - with major retailers already exploring device adoption.

WHY IT MATTERS

Project Solara answers “where do agents live outside the PC?” - giving developers a reference platform to build ambient, always-on agent experiences that escape the screen entirely and become woven into the physical environment.

Majorana 2

1,000× more reliable qubits · 20-second mean lifetime · targeting 2029

WHAT IT IS

Majorana 2 is Microsoft's next-generation topological quantum chip with a redesigned materials stack delivering 1,000× better qubit reliability, a mean qubit lifetime of 20 seconds (some instances lasting a full minute), and 1-microsecond gate operations.

WHAT IT DOES

These reliability advances cut Microsoft's roadmap to a commercially viable, fault-tolerant quantum computer in half - now targeting 2029 - with the redesigned materials stack (not just silicon) being the key engineering breakthrough.

WHY IT MATTERS

For developers, this signals quantum-classical hybrid computing is moving from a research curiosity to a credible near-decade timeline for enterprise production workloads - making "quantum-ready" architecture decisions increasingly relevant today.

Microsoft Discovery

AI agents applied to scientific research - now generally available

WHAT IT IS

Microsoft Discovery is now generally available - an agentic platform where scientists submit natural language research goals and receive AI-generated hypotheses, literature synthesis, and automated lab job submissions through a multi-step agentic reasoning loop.

WHAT IT DOES

Demonstrated live at Build with a chemistry researcher directing a full research workflow via three natural language prompts, Discovery also underpins a collaboration with Mayo Clinic to develop a frontier model specifically for healthcare.

WHY IT MATTERS

Discovery demonstrates that the same agentic stack powering enterprise automation can be applied to long-horizon scientific problems - AI agents are beginning to participate in science itself, not just accelerate writing about it.

Azure HorizonDB

Enterprise PostgreSQL for the AI era - OLTP + Vector + Memory in one engine

WHAT IT IS

Azure HorizonDB is a new enterprise-grade PostgreSQL engine combining OLTP, vector indexing, and agent memory storage in a single managed database - with vector as a first-class index type alongside B-tree and hash, not a bolt-on plugin.

WHAT IT DOES

Developers building agents no longer need separate relational and vector databases; HorizonDB handles transactional workloads and semantic retrieval in one engine with native Azure integration, removing the complexity of two-tier database architectures.

WHY IT MATTERS

Making vector a first-class index primitive on par with rows and keys reorders the AI application database stack - comparable to when NoSQL first made document storage a credible alternative to relational schemas.

Azure Cobalt 200 & AI Infrastructure

50% faster Arm VMs · 200 Gbps networking · Azure Linux 4.0

WHAT IT IS

Azure Cobalt 200 is Microsoft's new Arm-based VM series delivering 50% better performance for AI workloads, paired with Lasv5/Laosv5 VMs adding 35% more CPU, 138 TB local NVMe storage, and 200 Gbps network bandwidth.

WHAT IT DOES

The Build infrastructure lineup also includes Guest RDMA for Azure Boost (100 Gb/s low-latency interconnect for distributed AI and HPC), Confidential Live Migration for Intel TDX VMs, and Azure Linux 4.0 extending container hosting to general-purpose VMs.

WHY IT MATTERS

For developers running large-scale training, inference, or agent fleets on Azure, Cobalt 200 means significantly cheaper and faster cloud workloads with no code changes required - just choose the new instance type.

Surface RTX Spark Dev Box

1 petaflop of local AI · 128 GB memory · 120B parameter models on-device

WHAT IT IS

The Surface RTX Spark Dev Box is a compact developer workstation powered by NVIDIA's RTX Spark SoC, delivering 1 petaflop of AI performance with 128 GB of unified memory and support for running local models up to 120 billion parameters.


WHAT IT DOES

It enables developers to fine-tune models, run inference, test agentic pipelines, and prototype with state-of-the-art LLMs entirely on-device - no cloud round-trips, no data leaving the machine.

WHY IT MATTERS

For AI developers forced to choose between laptop convenience and cloud VM cluster power, RTX Spark Dev Box closes that gap - providing cloud-class local AI performance at a developer workstation price with data fully on-premises.

Copilot Pricing ...



Copilot Credits & The Agent Economy

Consumption-based pricing for autonomous agent work - \$0.01 per credit

WHAT IT IS

Microsoft introduced Copilot Credits as the new consumption meter for autonomous agent work - agents operate pay-as-you-go at \$0.01/credit or via prepaid credit packs, completely separated from per-user human Copilot licenses which remain unchanged.

WHAT IT DOES

This pricing model separates the cost of human-in-the-loop Copilot (per seat) from autonomous agent workloads (per credit), letting organizations scale agent automation without multiplying their per-seat license count.

WHY IT MATTERS

Copilot Credits create the first predictable, measurable unit of agent compute - making it possible to budget, optimize, and audit agentic automation at enterprise scale, treating AI agents as infrastructure cost rather than headcount.

Build 2026: The Agent Stack Is Here

MODELS

7 first-party MAI models · Aion on-device SLMs · Frontier Tuning

AGENTS

Scout Autopilot · OpenClaw open-source · Agent Framework 1.0 GA

CONTEXT

Work IQ · Web IQ · Foundry IQ - production grounding everywhere

GOVERNANCE

Agent 365 SDK GA · ACS open standard · MXC · Copilot Credits

SECURITY

MDASH 96.55% CyberGym · Defender + GitHub Code Security GA

GITHUB

Copilot App · Agent Merge · Copilot SDK GA in 6 languages

WINDOWS

Agent OS · WSL containers · Coreutils · Intelligent Terminal


HARDWARE

Project Solara · Majorana 2 · RTX Spark Dev Box · Cobalt 200 +50%

“The question is no longer whether agents will run at work. It’s whether your platform is ready to govern them.”

One more thing...

Thoughts on Software Development
from Mark Russinovich & Scott Hanselman



Generative AI as Seniority-Biased Technological Change: Evidence from U.S. Résumé and Job Posting Data

35 Pages • Posted: 8 Sep 2025

Guy Lichtinger


Harvard University Department of Economics

Seyed Mahdi Hosseini Maasoum

Harvard University

Date Written: August 31, 2025

Beginning in 2023Q1, **junior** employment in adopting firms **declined sharply** relative to non adopters, while **senior** employment **continued to rise**.



Since the widespread adoption of generative AI, **early-career workers** (age 22-25) in most AI-exposed occupations have experienced a **13 percent relative decline** in employment even after controlling for firm-level shocks.

In contrast, the employment for workers in less exposed fields and more experienced workers in the same occupations has remained stable or continued to grow.


Canaries in the Coal Mine? Six Facts about the Recent Employment Effects of Artificial Intelligence

Erik Brynjolfsson*

Bharat Chandar†

Ruyu Chen‡§¶

Headcount Over Time by Age Group
Software Developers (Normalized)



Software engineering is much more than knowing
a programming language or framework



[Popular](#)[Latest](#)[Newsletters](#)*The Atlantic*[Sign In](#)[Subscribe](#)

TECHNOLOGY

There's Never Been a Better Time to Study Computer Science

Even as AI progresses, coders aren't doomed.

By Lila Shroff





Illustration by The Atlantic. Source: Getty.

Thank You!

martina.grom@atwork.at · MVP Foundry and M365
toni.pohl@atwork.at · MVP Foundry and M365 Dev



//LOCALHOST:VIENNA

<https://globalai.community/chapters/vienna>

We are looking for sponsors and locations. Please get in touch if you know of anything!

